

Pong from Pixels

Deep RL Bootcamp

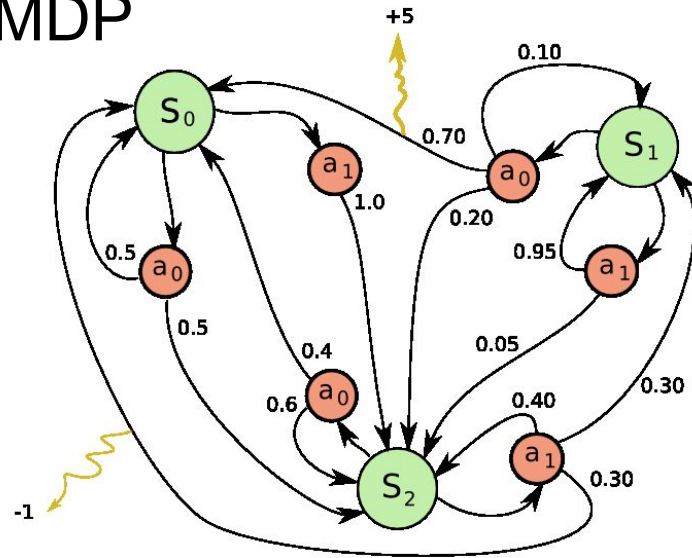
Andrej Karpathy, Aug 26, 2017

1


2

3

MDP



Policy network




Policy network

e.g.,


height width

[80 x 80]

array of



Policy network



height width


[80 x 80]

array

```
h = np.dot(W1, x) # compute hidden layer neuron activations
h[h<0] = 0 # ReLU nonlinearity: threshold at zero
logp = np.dot(W2, h) # compute log probability of going up
p = 1.0 / (1.0 + np.exp(-logp)) # sigmoid function (gives probability of going up)
```

Policy network

height width
[80 x 80]
array




E.g. 200 nodes in the hidden network, so:

$$[(80*80)*200 + 200] + [200*1 + 1] = \sim 1.3\text{M parameters}$$

Layer 1

Layer 2



Network does not see this. Network sees $80 \times 80 = 6,400$ numbers.
It gets a reward of +1 or -1, some of the time.
Q: How do we efficiently find a good setting of the 1.3M parameters?

Problem is easy if you want to be inefficient...

1. Repeat Forever:

2. Sample 1.3M random numbers
3. Run the policy for a while
4. If the performance is best so far, save it
5. Return the best policy

Problem is easy if you want to be inefficient...

1. **Repeat**
2. **Save**
3. **Run**
4. **If the**
5. **Return**



Problem is easy if you want to be inefficient...

1. **Repeat**
2. **Save**
3. **Run**
4. **If th**
5. **Return**



Policy Gradients




Suppose we had the training labels...
(we know what to do in any state)

(x1,UP)
(x2,DOWN)
(x3,UP)
...

Suppose we had the training labels...
(we know what to do in any state)

(x1,UP)
(x2,DOWN)
(x3,UP)
...




Suppose we had the training labels...
(we know what to do in any state)



(x1,UP)
(x2,DOWN)
(x3,UP)
...

maximize:

$$\sum_i \log p(y_i | x_i)$$



Except, we don't have labels...



Should we go UP or DOWN?


Except, we don't have labels...




“Try a bunch of stuff and see what happens. Do more of the stuff that worked in the future.”

-RL

Let's just act according to our current policy...

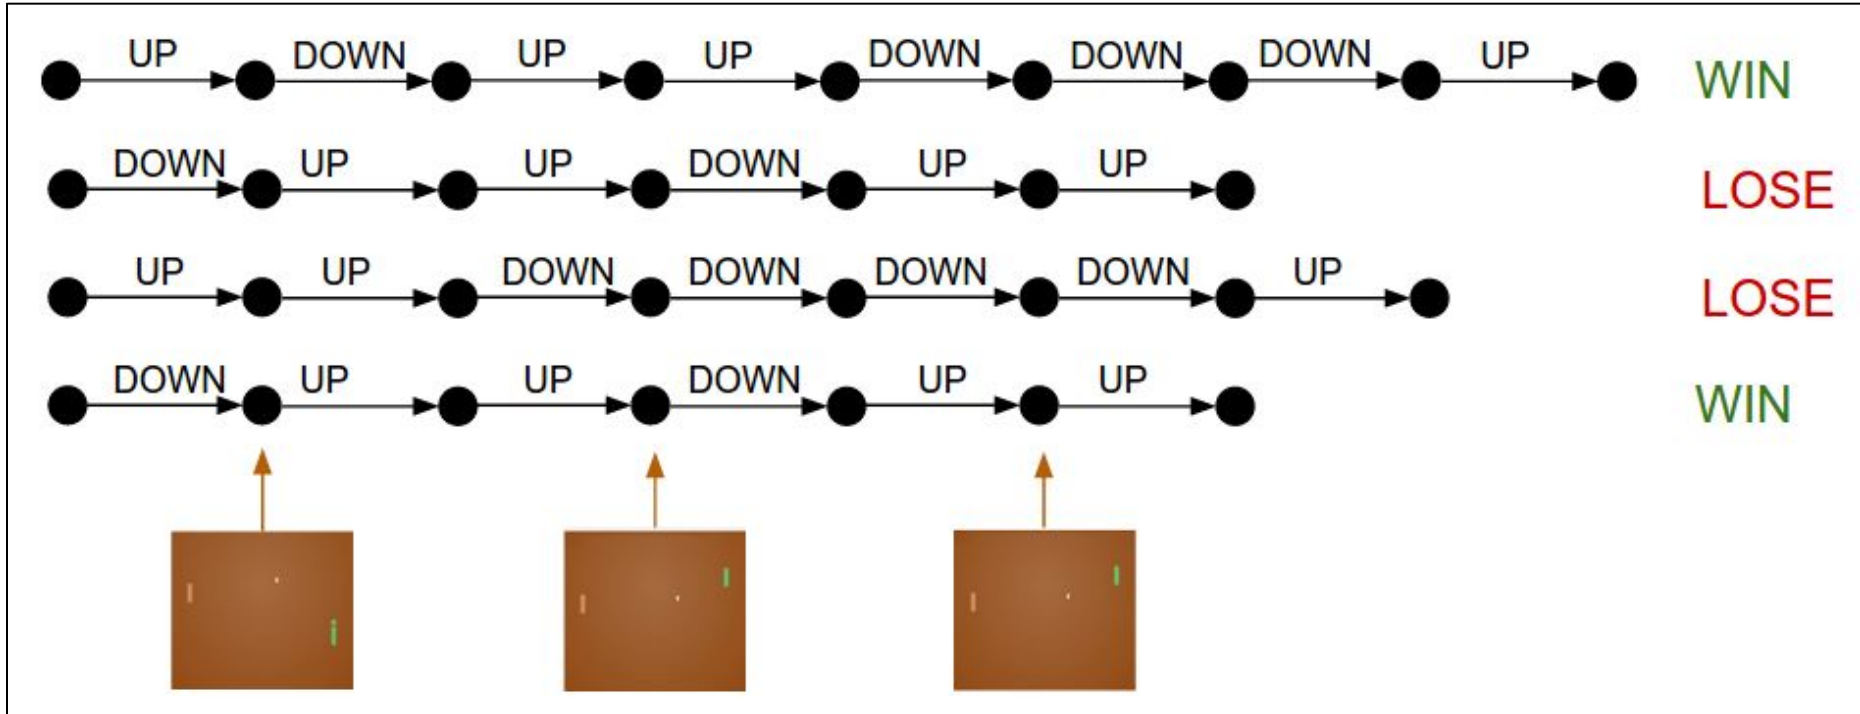


Rollout the policy
and collect an
episode

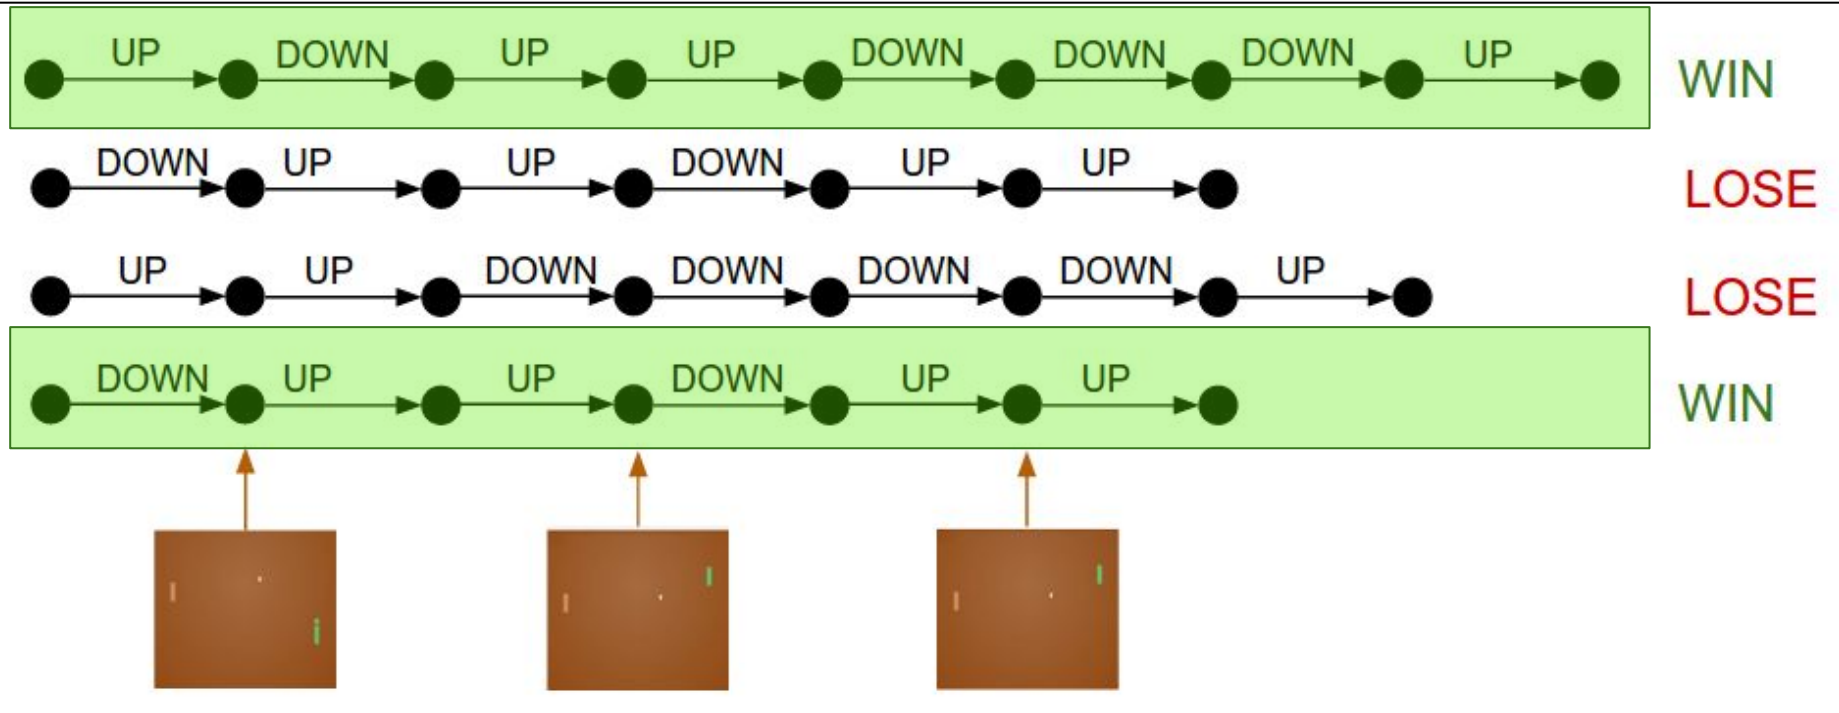


Collect many rollouts...

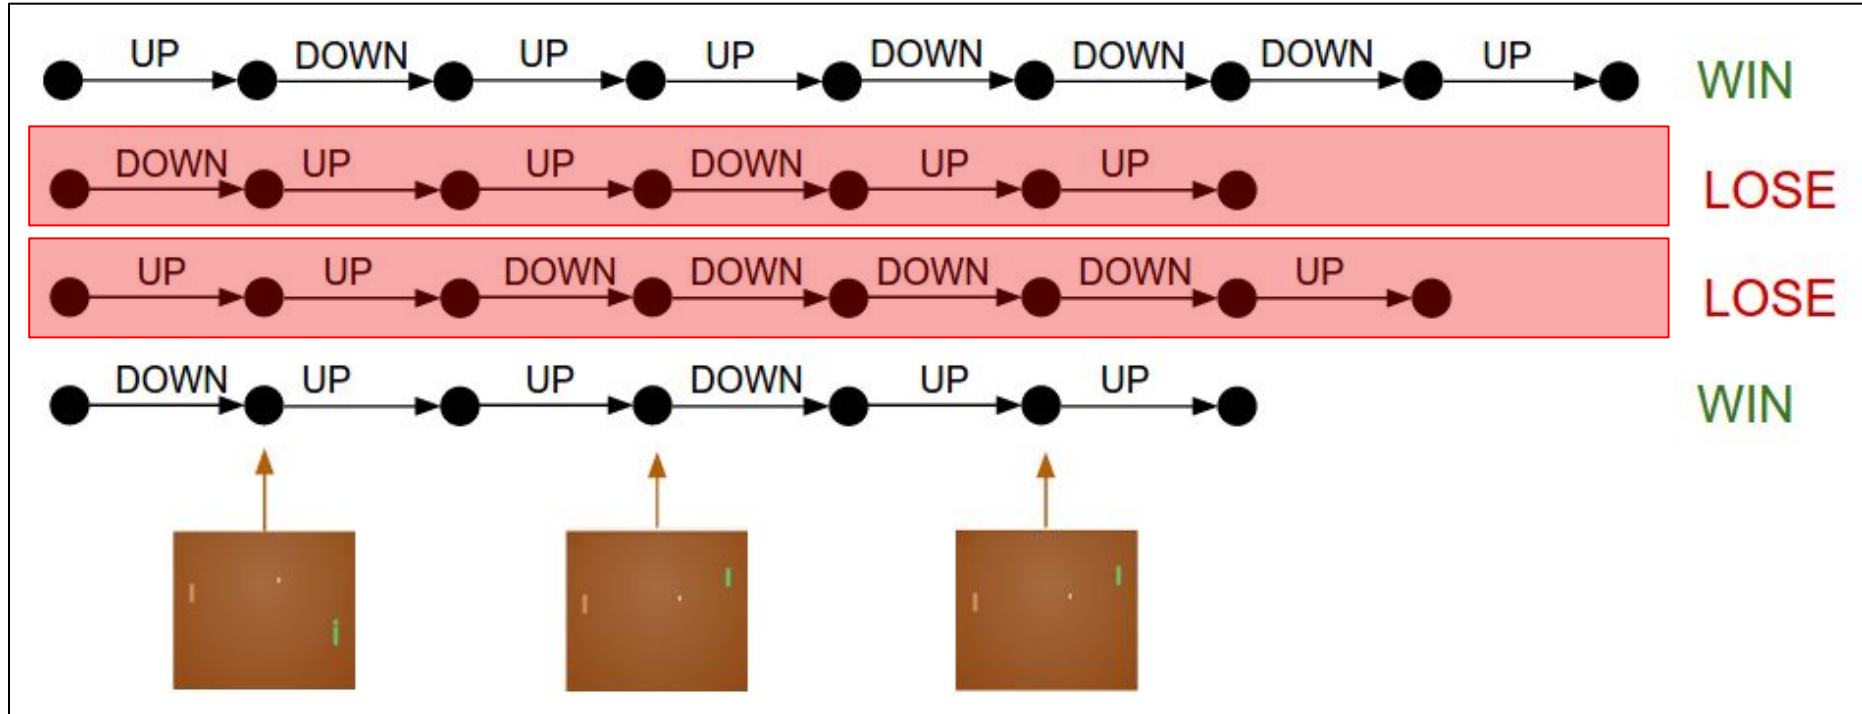
4 rollouts:



Not sure whatever we did here, but
apparently it was good.



Not sure whatever we did here, but it was bad.

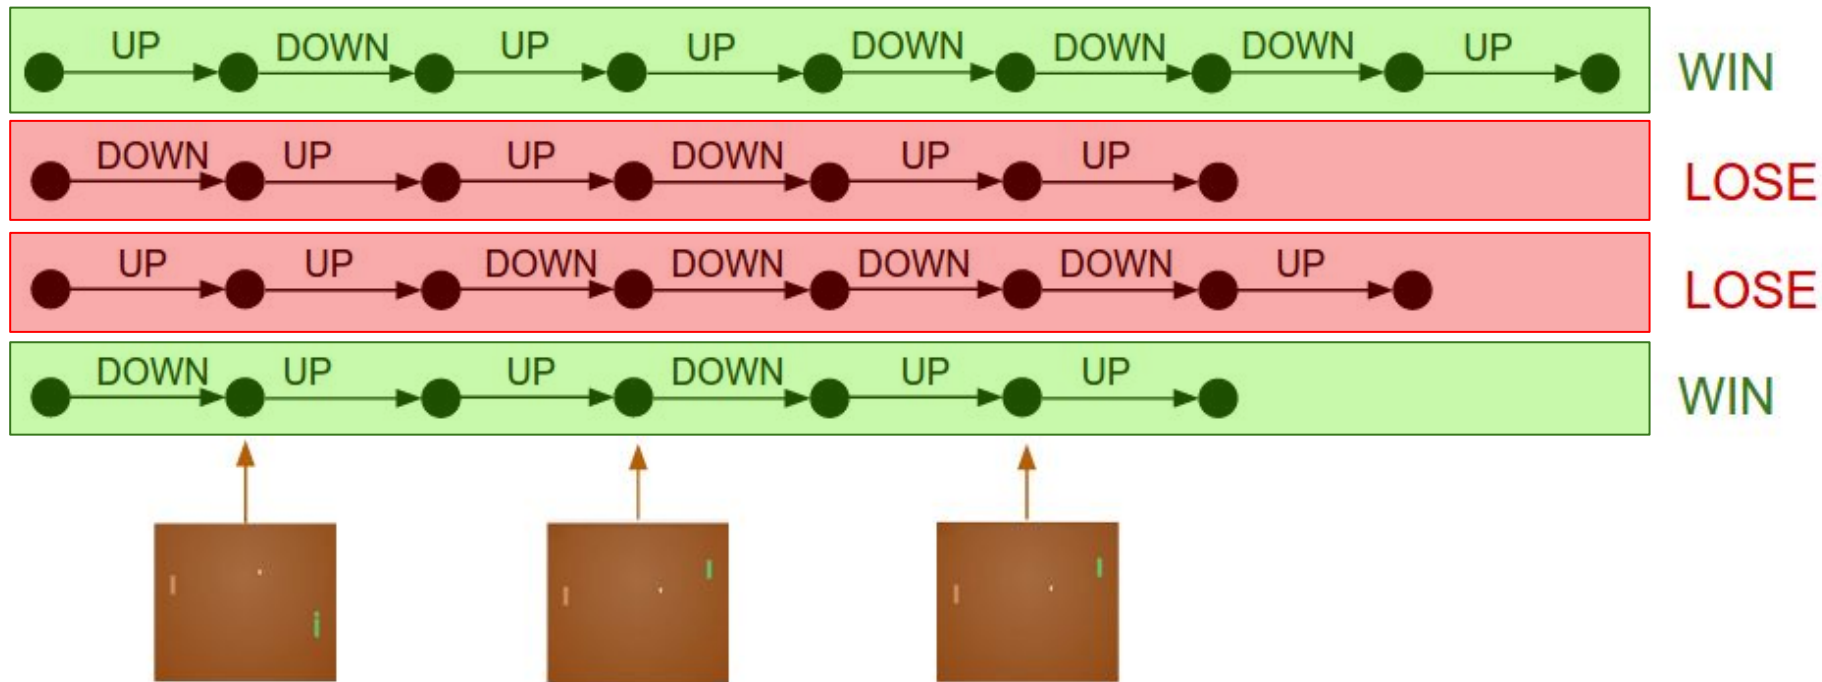


Pretend every action we took here was the correct label.

maximize: $\log p(y_i | x_i)$

Pretend every action we took here was the wrong label.

maximize: $(-1) * \log p(y_i | x_i)$



Supervised Learning

maximize:

$$\sum_i \log p(y_i | x_i)$$

For images x_i and their labels y_i .

Supervised Learning

maximize:

$$\sum_i \log p(y_i | x_i)$$

For images x_i and their labels y_i .

Reinforcement Learning

Supervised Learning

maximize:

$$\sum_i \log p(y_i | x_i)$$

For images x_i and their labels y_i .

Reinforcement Learning

1) we have no labels so we sample:

$$y_i \sim p(\cdot | x_i)$$

Supervised Learning

maximize:

$$\sum_i \log p(y_i | x_i)$$

For images x_i and their labels y_i .

Reinforcement Learning

1) we have no labels so we sample:

$$y_i \sim p(\cdot | x_i)$$

2) once we collect a batch of rollouts:
maximize:

$$\sum_i A_i * \log p(y_i | x_i)$$

Supervised Learning

maximize:

$$\sum_i \log p(y_i | x_i)$$

For images x_i and their labels y_i .

Reinforcement Learning

1) we have no labels so we sample:

$$y_i \sim p(\cdot | x_i)$$

2) once we collect a batch of rollouts:
maximize:

$$\sum_i A_i * \log p(y_i | x_i)$$

We call this the **advantage**, it's a number, like +1.0 or -1.0 based on how this action eventually turned out.

Supervised Learning

maximize:

$$\sum_i \log p(y_i | x_i)$$

For images x_i and their labels y_i .

Reinforcement Learning

1) we have no labels so we sample:

$$y_i \sim p(\cdot | x_i)$$

2) once we collect a batch of rollouts:
maximize:

$$\sum_i A_i * \log p(y_i | x_i)$$

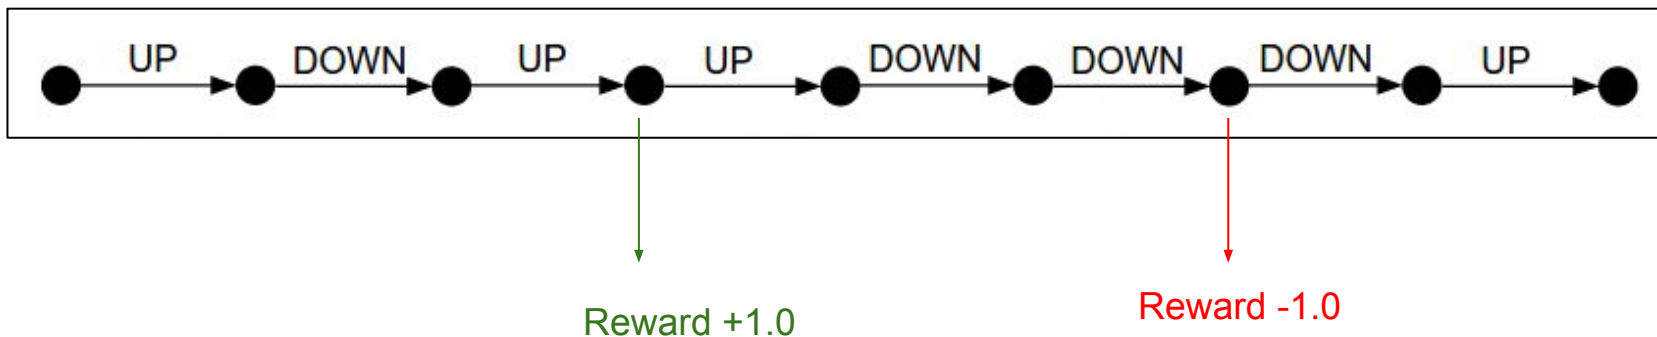

+ve advantage will make that action more likely in the future, for that state.

-ve advantage will make that action less likely in the future, for that state.



Discounting

Blame each action assuming that its effects have exponentially decaying impact into the future.




Discounting

Blame each action assuming that its effects have exponentially decaying impact into the future.

Discounted rewards

$$\sum_i A_i * \log p(y_i | x_i)$$


0.21 0.24 0.27 -0.81 -0.9 -1 0 0



Reward +1.0

Reward -1.0

$\gamma = 0.9$



**YEAH, IF YOU COULD JUST SHOW ME
SOME CODE**



THAT'D BE GREAT

130 line gist, numpy as the only dependency.

<https://gist.github.com/karpathy/a4166c7fe253700972fcbc77e4ea32c5>

```

64 env = gym.make("Pong-v0")
65 observation = env.reset()
66 prev_x = None # used in computing the difference frame
67 xs,hs,dlogps,drs = [],[],[],[]
68 running_reward = None
69 reward_sum = 0
70 episode_number = 0
71 while True:
72     if render: env.render()
73
74     # preprocess the observation, set input to network to be difference image
75     cur_x = prepro(observation)
76     x = cur_x - prev_x if prev_x is not None else np.zeros(D)
77     prev_x = cur_x
78
79     # forward the policy network and sample an action from the returned probability
80     aprobs, h = policy_forward(x)
81     action = 2 if np.random.uniform() < aprobs else 3 # roll the dice!
82
83     # record various intermediates (needed later for backprop)
84     xs.append(x) # observation
85     hs.append(h) # hidden state
86     y = 1 if action == 2 else 0 # a "fake label"
87     dlogps.append(y - aprobs) # grad that encourages the action that was taken to be taken (see http://cs231n.github.io/neural-networks-2/#loss)
88
89     # step the environment and get new measurements
90     observation, reward, done, info = env.step(action)
91     reward_sum += reward
92
93     drs.append(reward) # record reward (has to be done after we call step() to get reward for previous action)
94
95     if done: # an episode finished
96         episode_number += 1
97
98         # stack together all inputs, hidden states, action gradients, and rewards for this episode
99
100         eps = np.vstack(xs)
101         eph = np.vstack(hs)
102         epdlogp = np.vstack(dlogps)
103         epr = np.vstack(drs)
104         xs,hs,dlogps,drs = [],[],[],[] # reset array memory
105
106         # compute the discounted reward backwards through time
107         discounted_epr = discount_rewards(epr)
108         # standardize the rewards to be unit normal (helps control the gradient estimator variance)
109         discounted_epr -= np.mean(discounted_epr)
110         discounted_epr /= np.std(discounted_epr)
111
112         epdlogp *= discounted_epr # modulate the gradient with advantage (PG magic happens right here.)
113         grad = policy_backward(eph, epdlogp)
114         for k in model: grad_buffer[k] += grad[k] # accumulate grad over batch
115
116         # perform rmsprop parameter update every batch_size episodes
117         if episode_number % batch_size == 0:
118             for k,v in model.iteritems():
119                 g = grad_buffer[k] # gradient
120                 rmsprop_cache[k] = decay_rate * rmsprop_cache[k] + (1 - decay_rate) * g**2
121                 model[k] += learning_rate * g / (np.sqrt(rmsprop_cache[k] + 1e-5))
122                 grad_buffer[k] = np.zeros_like(v) # reset batch gradient buffer
123
124         # boring book-keeping
125         running_reward = reward_sum if running_reward is None else running_reward * 0.99 + reward_sum * 0.01
126         print 'resetting env. episode reward total was %f. running mean: %f' % (reward_sum, running_reward)
127         if episode_number % 100 == 0: pickle.dump(model, open('save.p', 'wb'))
128         reward_sum = 0
129         observation = env.reset() # reset env
130         prev_x = None
131
132     if reward != 0: # Pong has either +1 or -1 reward exactly when game ends.
133         print ('ep %d: game finished, reward: %f' % (episode_number, reward)) + (' ' if reward == -1 else '!!!!!!!')

```

```

env = gym.make("Pong-v0")
observation = env.reset()

prev_x = None # used in computing the difference frame
xs,hs,dlogps,drs = [],[],[],[]

running_reward = None
reward_sum = 0
episode_number = 0

while True:
    if render: env.render()

```

Nothing too scary over here.

We use OpenAI Gym.
And start the main training loop.

```

64 env = gym.make("Pong-v0")
65 observation = env.reset()
66 prev_x = None # used in computing the difference frame
67 xs,hs,dlogps,drs = [],[],[],[]
68 running_reward = None
69 reward_sum = 0
70 episode_number = 0
71 while True:
72     if render: env.render()
73
74     # preprocess the observation, set input to network to be difference image
75     cur_x = prepro(observation)
76     x = cur_x - prev_x if prev_x is not None else np.zeros(D)
77     prev_x = cur_x
78
79     # forward the policy network and sample an action from the returned probability
80     aprobs, h = policy_forward(x)
81     action = 2 if np.random.uniform() < aprobs else 3 # roll the dice!
82
83     # record various intermediates (needed later for backprop)
84     xs.append(x) # observation
85     hs.append(h) # hidden state
86     y = 1 if action == 2 else 0 # a "fake label"
87     dlogps.append(y - aprobs) # grad that encourages the action that was taken to be taken (see http://cs231n.github.io/neural-networks-2/#loss)
88
89     # step the environment and get new measurements
90     observation, reward, done, info = env.step(action)
91     reward_sum += reward
92
93     drs.append(reward) # record reward (has to be done after we call step() to get reward for previous action)
94
95     if done: # an episode finished
96         episode_number += 1
97
98         # stack together all inputs, hidden states, action gradients, and rewards for this episode
99         eps = np.vstack(xs)
100         eph = np.vstack(hs)
101         epdlogps = np.vstack(dlogps)
102         epr = np.vstack(drs)
103         xs,hs,dlogps,drs = [],[],[],[] # reset array memory
104
105         # compute the discounted reward backwards through time
106         discounted_epr = discount_rewards(epr)
107         # standardize the rewards to be unit normal (helps control the gradient estimator variance)
108         discounted_epr -= np.mean(discounted_epr)
109         discounted_epr /= np.std(discounted_epr)
110
111         epdlogps *= discounted_epr # modulate the gradient with advantage (PG magic happens right here.)
112         grad = policy_backward(eph, epdlogps)
113         for k in model: grad_buffer[k] += grad[k] # accumulate grad over batch
114
115         # perform rmsprop parameter update every batch_size episodes
116         if episode_number % batch_size == 0:
117             for k,v in model.iteritems():
118                 g = grad_buffer[k] # gradient
119                 rmsprop_cache[k] = decay_rate * rmsprop_cache[k] + (1 - decay_rate) * g**2
120                 model[k] += learning_rate * g / (np.sqrt(rmsprop_cache[k]) + 1e-5)
121                 grad_buffer[k] = np.zeros_like(v) # reset batch gradient buffer
122
123         # boring book-keeping
124         running_reward = reward_sum if running_reward is None else running_reward * 0.99 + reward_sum * 0.01
125         print 'resetting env. episode reward total was %f. running mean: %f' % (reward_sum, running_reward)
126         if episode_number % 100 == 0: pickle.dump(model, open('save.p', 'wb'))
127         reward_sum = 0
128         observation = env.reset() # reset env
129         prev_x = None
130
131     if reward != 0: # Pong has either +1 or -1 reward exactly when game ends.
132         print ('ep %d: game finished, reward: %f' % (episode_number, reward)) + (' ' if reward == -1 else '!!!!!!!')

```

```

# preprocess the observation, set input to network to be difference image
cur_x = prepro(observation)
x = cur_x - prev_x if prev_x is not None else np.zeros(D)
prev_x = cur_x

```

```

def prepro(I):
    """ prepro 210x160x3 uint8 frame into 6400 (80x80) 1D float vector """
    I = I[35:195] # crop
    I = I[:,2,:2,0] # downsample by factor of 2
    I[I == 144] = 0 # erase background (background type 1)
    I[I == 109] = 0 # erase background (background type 2)
    I[I != 0] = 1 # everything else (paddles, ball) just set to 1
    return I.astype(np.float).ravel()

```

Get the current image and preprocess it.

```

64 env = gym.make("Pong-v0")
65 observation = env.reset()
66 prev_x = None # used in computing the difference frame
67 xs,hs,dlogps,drs = [],[],[],[]
68 running_reward = None
69 reward_sum = 0
70 episode_number = 0
71 while True:
72     if render: env.render()
73
74     # preprocess the observation, set input to network to be difference image
75     cur_x = prepro(observation)
76     x = cur_x - prev_x if prev_x is not None else np.zeros(D)
77     prev_x = cur_x
78
79     # forward the policy network and sample an action from the returned probability
80     aprob, h = policy_forward(x)
81     action = 2 if np.random.uniform() < aprob else 3 # roll the dice!
82
83     # record various intermediates (needed later for backprop)
84     xs.append(x) # observation
85     hs.append(h) # hidden state
86     y = 1 if action == 2 else 0 # a "fake label"
87     dlogps.append(y - aprob) # grad that encourages the action that was taken to be taken (see http://cs231n.github.io/neural-networks-2/#loss)
88
89     # step the environment and get new measurements
90     observation, reward, done, info = env.step(action)
91     reward_sum += reward
92
93     drs.append(reward) # record reward (has to be done after we call step() to get reward for previous action)
94
95     if done: # an episode finished
96         episode_number += 1
97
98         # stack together all inputs, hidden states, action gradients, and rewards for this episode
99         eps = np.vstack(xs)
100         eph = np.vstack(hs)
101         epdlogp = np.vstack(dlogps)
102         epr = np.vstack(drs)
103         xs,hs,dlogps,drs = [],[],[],[] # reset array memory
104
105         # compute the discounted reward backwards through time
106         discounted_epr = discount_rewards(epr)
107         # standardize the rewards to be unit normal (helps control the gradient estimator variance)
108         discounted_epr -= np.mean(discounted_epr)
109         discounted_epr /= np.std(discounted_epr)
110
111         epdlogp *= discounted_epr # modulate the gradient with advantage (PG magic happens right here.)
112         grad = policy_backward(eph, epdlogp)
113         for k in model: grad_buffer[k] += grad[k] # accumulate grad over batch
114
115         # perform rmsprop parameter update every batch_size episodes
116         if episode_number % batch_size == 0:
117             for k,v in model.iteritems():
118                 g = grad_buffer[k] # gradient
119                 rmsprop_cache[k] = decay_rate * rmsprop_cache[k] + (1 - decay_rate) * g**2
120                 model[k] += learning_rate * g / (np.sqrt(rmsprop_cache[k]) + 1e-5)
121                 grad_buffer[k] = np.zeros_like(v) # reset batch gradient buffer
122
123         # boring book-keeping
124         running_reward = reward_sum if running_reward is None else running_reward * 0.99 + reward_sum * 0.01
125         print 'resetting env. episode reward total was %f. running mean: %f' % (reward_sum, running_reward)
126         if episode_number % 100 == 0: pickle.dump(model, open('save.p', 'wb'))
127         reward_sum = 0
128         observation = env.reset() # reset env
129         prev_x = None
130
131     if reward != 0: # Pong has either +1 or -1 reward exactly when game ends.
132         print ('ep %d: game finished, reward: %f' % (episode_number, reward)) + (' ' if reward == -1 else '!!!!!!!')

```

```

# forward the policy network and sample an action from the returned probability
aprob, h = policy_forward(x)
action = 2 if np.random.uniform() < aprob else 3 # roll the dice!

```

```

def policy_forward(x):
    h = np.dot(model['W1'], x)
    h[h<0] = 0 # ReLU nonlinearity
    logp = np.dot(model['W2'], h)
    p = sigmoid(logp)
    return p, h # return probability of taking action 2, and hidden state

```

```

def sigmoid(x):
    return 1.0 / (1.0 + np.exp(-x)) # sigmoid "squashing" function to interval [0,1]

```

```

64 env = gym.make("Pong-v0")
65 observation = env.reset()
66 prev_x = None # used in computing the difference frame
67 xs,hs,dlogps,drs = [],[],[],[]
68 running_reward = None
69 reward_sum = 0
70 episode_number = 0
71 while True:
72     if render: env.render()
73
74     # preprocess the observation, set input to network to be difference image
75     cur_x = prepro(observation)
76     x = cur_x - prev_x if prev_x is not None else np.zeros(D)
77     prev_x = cur_x
78
79     # forward the policy network and sample an action from the returned probability
80     aprob, h = policy_forward(x)
81     action = 2 if np.random.uniform() < aprob else 3 # roll the dice!
82
83     # record various intermediates (needed later for backprop)
84     xs.append(x) # observation
85     hs.append(h) # hidden state
86     y = 1 if action == 2 else 0 # a "fake label"
87     dlogps.append(y - aprob) # grad that encourages the action that was taken to be taken (see http://cs231n.github.io/neural-networks-2/#loss)
88
89     # step the environment and get new measurements
90     observation, reward, done, info = env.step(action)
91     reward_sum += reward
92
93     drs.append(reward) # record reward (has to be done after we call step() to get reward for previous action)
94
95     if done: # an episode finished
96         episode_number += 1
97
98         # stack together all inputs, hidden states, action gradients, and rewards for this episode
99         eps = np.vstack(xs)
100         eph = np.vstack(hs)
101         epdlogp = np.vstack(dlogps)
102         epr = np.vstack(drs)
103         xs,hs,dlogps,drs = [],[],[],[] # reset array memory
104
105         # compute the discounted reward backwards through time
106         discounted_epr = discount_rewards(epr)
107         # standardize the rewards to be unit normal (helps control the gradient estimator variance)
108         discounted_epr -= np.mean(discounted_epr)
109         discounted_epr /= np.std(discounted_epr)
110
111         epdlogp *= discounted_epr # modulate the gradient with advantage (PG magic happens right here.)
112         grad = policy_backward(eph, epdlogp)
113         for k in model: grad_buffer[k] += grad[k] # accumulate grad over batch
114
115         # perform rmsprop parameter update every batch_size episodes
116         if episode_number % batch_size == 0:
117             for k,v in model.iteritems():
118                 g = grad_buffer[k] # gradient
119                 rmsprop_cache[k] = decay_rate * rmsprop_cache[k] + (1 - decay_rate) * g**2
120                 model[k] += learning_rate * g / (np.sqrt(rmsprop_cache[k] + 1e-5))
121                 grad_buffer[k] = np.zeros_like(v) # reset batch gradient buffer
122
123         # boring book-keeping
124         running_reward = reward_sum if running_reward is None else running_reward * 0.99 + reward_sum * 0.01
125         print 'resetting env. episode reward total was %f. running mean: %f' % (reward_sum, running_reward)
126         if episode_number % 100 == 0: pickle.dump(model, open('save.p', 'wb'))
127         reward_sum = 0
128         observation = env.reset() # reset env
129         prev_x = None
130
131     if reward != 0: # Pong has either +1 or -1 reward exactly when game ends.
132         print ('ep %d: game finished, reward: %f' % (episode_number, reward)) + (' ' if reward == -1 else '!!!!!!!')

```

```

# record various intermediates (needed later for backprop)
xs.append(x) # observation
hs.append(h) # hidden state
y = 1 if action == 2 else 0 # a "fake label"
dlogps.append(y - aprob) # grad that encourages the action that was taken to be taken

```

Bookkeeping so that we can do backpropagation later. If you were to use PyTorch or something, this would not be needed.


```

64 env = gym.make("Pong-v0")
65 observation = env.reset()
66 prev_x = None # used in computing the difference frame
67 xs,hs,dlogps,drs = [],[],[],[]
68 running_reward = None
69 reward_sum = 0
70 episode_number = 0
71 while True:
72     if render: env.render()
73
74     # preprocess the observation, set input to network to be difference image
75     cur_x = prepro(observation)
76     x = cur_x - prev_x if prev_x is not None else np.zeros(D)
77     prev_x = cur_x
78
79     # forward the policy network and sample an action from the returned probability
80     aprobs, h = policy_forward(x)
81     action = 2 if np.random.uniform() < aprobs else 3 # roll the dice!
82
83     # record various intermediates (needed later for backprop)
84     xs.append(x) # observation
85     hs.append(h) # hidden state
86     y = 1 if action == 2 else 0 # a "fake label"
87     dlogps.append(y - aprobs) # grad that encourages the action that was taken (see http://cs231n.github.io/neural-networks-2/#loss)
88
89     # step the environment and get new measurements
90     observation, reward, done, info = env.step(action)
91     reward_sum += reward
92
93     ddrs.append(reward) # record reward (has to be done after we call step() to get reward for previous action)
94
95     if done: # an episode finished
96         episode_number += 1
97
98         # stack together all inputs, hidden states, action gradients, and rewards for this episode
99         eps = np.vstack(xs)
100         eph = np.vstack(hs)
101         epdlogps = np.vstack(dlogps)
102         epr = np.vstack(drs)
103         xs,hs,dlogps,drs = [],[],[],[] # reset array memory
104
105         # compute the discounted reward backwards through time
106         discounted_epr = discount_rewards(epr)
107         # standardize the rewards to be unit normal (helps control the gradient estimator variance)
108         discounted_epr -= np.mean(discounted_epr)
109         discounted_epr /= np.std(discounted_epr)
110
111         epdlogps *= discounted_epr # modulate the gradient with advantage (PG magic happens right here.)
112         grad = policy_backward(eph, epdlogps)
113         for k in model: grad_buffer[k] += grad[k] # accumulate grad over batch
114
115         # perform rmsprop parameter update every batch_size episodes
116         if episode_number % batch_size == 0:
117             for k,v in model.iteritems():
118                 g = grad_buffer[k] # gradient
119                 rmsprop_cache[k] = decay_rate * rmsprop_cache[k] + (1 - decay_rate) * g**2
120                 model[k] += learning_rate * g / (np.sqrt(rmsprop_cache[k]) + 1e-5)
121                 grad_buffer[k] = np.zeros_like(v) # reset batch gradient buffer
122
123         # boring book-keeping
124         running_reward = reward_sum if running_reward is None else running_reward * 0.99 + reward_sum * 0.01
125         print 'resetting env. episode reward total was %f. running mean: %f' % (reward_sum, running_reward)
126         if episode_number % 100 == 0: pickle.dump(model, open('save.p', 'wb'))
127         reward_sum = 0
128         observation = env.reset() # reset env
129         prev_x = None
130
131     if reward != 0: # Pong has either +1 or -1 reward exactly when game ends.
132         print ('ep %d: game finished, reward: %f' % (episode_number, reward)) + (' ' if reward == -1 else '!!!!!!!')

```

```

# record various intermediates (needed later for backprop)
xs.append(x) # observation
hs.append(h) # hidden state
y = 1 if action == 2 else 0 # a "fake label"
dlogps.append(y - aprobs) # grad that encourages the action that was taken to be taken

```

A small piece of backprop:

Derivative of the [log probability of the taken action given this image] with respect to the [output of the network (before sigmoid)]

recall: loss:

$$\sum_i A_i * \log p(y_i | x_i)$$

$$s = W_2 f(W_1 x)$$

$$p = 1 / (1 + e^{-s})$$

$$y \sim p$$

```

64 env = gym.make("Pong-v0")
65 observation = env.reset()
66 prev_x = None # used in computing the difference frame
67 xs,hs,dlogps,drs = [],[],[],[]
68 running_reward = None
69 reward_sum = 0
70 episode_number = 0
71 while True:
72     if render: env.render()
73
74     # preprocess the observation, set input to network to be difference image
75     cur_x = prepro(observation)
76     x = cur_x - prev_x if prev_x is not None else np.zeros(D)
77     prev_x = cur_x
78
79     # forward the policy network and sample an action from the returned probability
80     aprobs, h = policy_forward(x)
81     action = 2 if np.random.uniform() < aprobs else 3 # roll the dice!
82
83     # record various intermediates (needed later for backprop)
84     xs.append(x) # observation
85     hs.append(h) # hidden state
86     y = 1 if action == 2 else 0 # a "fake label"
87     dlogps.append(y - aprobs) # grad that encourages the action that was taken (see http://cs231n.github.io/neural-networks-2/#loss)
88
89     # step the environment and get new measurements
90     observation, reward, done, info = env.step(action)
91     reward_sum += reward
92
93     drs.append(reward) # record reward (has to be done after we call step() to get reward for previous action)
94
95     if done: # an episode finished
96         episode_number += 1
97
98         # stack together all inputs, hidden states, action gradients, and rewards for this episode
99         eps = np.vstack(xs)
100         eph = np.vstack(hs)
101         epdlogps = np.vstack(dlogps)
102         epr = np.vstack(drs)
103         xs,hs,dlogps,drs = [],[],[],[] # reset array memory
104
105         # compute the discounted reward backwards through time
106         discounted_epr = discount_rewards(epr)
107         # standardize the rewards to be unit normal (helps control the gradient estimator variance)
108         discounted_epr -= np.mean(discounted_epr)
109         discounted_epr /= np.std(discounted_epr)
110
111         epdlogps *= discounted_epr # modulate the gradient with advantage (PG magic happens right here.)
112         grad = policy_backward(eph, epdlogps)
113         for k in model: grad_buffer[k] += grad[k] # accumulate grad over batch
114
115         # perform rmsprop parameter update every batch_size episodes
116         if episode_number % batch_size == 0:
117             for k,v in model.iteritems():
118                 g = grad_buffer[k] # gradient
119                 rmsprop_cache[k] = decay_rate * rmsprop_cache[k] + (1 - decay_rate) * g**2
120                 model[k] += learning_rate * g / (np.sqrt(rmsprop_cache[k]) + 1e-5)
121                 grad_buffer[k] = np.zeros_like(v) # reset batch gradient buffer
122
123         # boring book-keeping
124         running_reward = reward_sum if running_reward is None else running_reward * 0.99 + reward_sum * 0.01
125         print 'resetting env. episode reward total was %f. running mean: %f' % (reward_sum, running_reward)
126         if episode_number % 100 == 0: pickle.dump(model, open('save.p', 'wb'))
127         reward_sum = 0
128         observation = env.reset() # reset env
129         prev_x = None
130
131     if reward != 0: # Pong has either +1 or -1 reward exactly when game ends.
132         print ('ep %d: game finished, reward: %f' % (episode_number, reward)) + (' ' if reward == -1 else '!!!!!!!')

```

```

# record various intermediates (needed later for backprop)
xs.append(x) # observation
hs.append(h) # hidden state
y = 1 if action == 2 else 0 # a "fake label"
dlogps.append(y - aprobs) # grad that encourages the action that was taken to be taken

```

A small piece of backprop:

Derivative of the [log probability of the taken action given this image] with respect to the [output of the network (before sigmoid)]

recall: loss:

$$\sum_i A_i * \log p(y_i | x_i)$$

$$s = W_2 f(W_1 x)$$

$$p = 1 / (1 + e^{-s})$$

$$y \sim p$$

if $y = 1, L = \log p, dL/ds = 1 - p$

if $y = 0, L = \log(1 - p), dL/ds = -p$

$$L = y \log(p) + (1 - y) \log(1 - p)$$

$$dL/ds = y - p$$

More compact:


```

64 env = gym.make("Pong-v0")
65 observation = env.reset()
66 prev_x = None # used in computing the difference frame
67 xs,hs,dlogps,drs = [],[],[],[]
68 running_reward = None
69 reward_sum = 0
70 episode_number = 0
71 while True:
72     if render: env.render()
73
74     # preprocess the observation, set input to network to be difference image
75     cur_x = prepro(observation)
76     x = cur_x - prev_x if prev_x is not None else np.zeros(D)
77     prev_x = cur_x
78
79     # forward the policy network and sample an action from the returned probability
80     aprob, h = policy_forward(x)
81     action = 2 if np.random.uniform() < aprob else 3 # roll the dice!
82
83     # record various intermediates (needed later for backprop)
84     xs.append(x) # observation
85     hs.append(h) # hidden state
86     y = 1 if action == 2 else 0 # a "fake label"
87     dlogps.append(y - aprob) # grad that encourages the action that was taken to be taken (see http://cs231n.github.io/neural-networks-2/#los
88
89     # step the environment and get new measurements
90     observation, reward, done, info = env.step(action)
91     reward_sum += reward
92
93     drs.append(reward) # record reward (has to be done after we call step() to get reward for previous action)
94
95     if done: # an episode finished
96         episode_number += 1
97
98         # stack together all inputs, hidden states, action gradients, and rewards for this episode
99
100         eps = np.vstack(xs)
101         eph = np.vstack(hs)
102         epdlogp = np.vstack(dlogps)
103         epr = np.vstack(drs)
104         xs,hs,dlogps,drs = [],[],[],[] # reset array memory
105
106         # compute the discounted reward backwards through time
107         discounted_epr = discount_rewards(epr)
108         # standardize the rewards to be unit normal (helps control the gradient estimator variance)
109         discounted_epr -= np.mean(discounted_epr)
110         discounted_epr /= np.std(discounted_epr)
111
112         epdlogp *= discounted_epr # modulate the gradient with advantage (PG magic happens right here.)
113         grad = policy_backward(eph, epdlogp)
114         for k in model: grad_buffer[k] += grad[k] # accumulate grad over batch
115
116         # perform rmsprop parameter update every batch_size episodes
117         if episode_number % batch_size == 0:
118             for k,v in model.iteritems():
119                 g = grad_buffer[k] # gradient
120                 rmsprop_cache[k] = decay_rate * rmsprop_cache[k] + (1 - decay_rate) * g**2
121                 model[k] += learning_rate * g / (np.sqrt(rmsprop_cache[k] + 1e-5))
122                 grad_buffer[k] = np.zeros_like(v) # reset batch gradient buffer
123
124         # boring book-keeping
125         running_reward = reward_sum if running_reward is None else running_reward * 0.99 + reward_sum * 0.01
126         print 'resetting env. episode reward total was %f. running mean: %f' % (reward_sum, running_reward)
127         if episode_number % 100 == 0: pickle.dump(model, open('save.p', 'wb'))
128         reward_sum = 0
129         observation = env.reset() # reset env
130         prev_x = None
131
132     if reward != 0: # Pong has either +1 or -1 reward exactly when game ends.
133         print ('ep %d: game finished, reward: %f' % (episode_number, reward)) + (' ' if reward == -1 else '!!!!!!!')

```

```

# step the environment and get new measurements
observation, reward, done, info = env.step(action)
reward_sum += reward

drs.append(reward) # record reward (has to be done after we call step() to get reward for previous action)

```

Step the environment

(execute the action, get new state and record the reward)

```

64 env = gym.make("Pong-v0")
65 observation = env.reset()
66 prev_x = None # used in computing the difference frame
67 xs,hs,dlogps,drs = [],[],[],[]
68 running_reward = None
69 reward_sum = 0
70 episode_number = 0
71 while True:
72     if render: env.render()
73
74     # preprocess the observation, set input to network to be difference image
75     cur_x = prepro(observation)
76     x = cur_x - prev_x if prev_x is not None else np.zeros(D)
77     prev_x = cur_x
78
79     # forward the policy network and sample an action from the returned probability
80     aprobs, h = policy_forward(x)
81     action = 2 if np.random.uniform() < aprobs else 3 # roll the dice!
82
83     # record various intermediates (needed later for backprop)
84     xs.append(x) # observation
85     hs.append(h) # hidden state
86     y = 1 if action == 2 else 0 # a "fake label"
87     dlogps.append(y - aprobs) # grad that encourages the action that was taken to be taken (see http://cs231n.github.io/neural-networks-2/#los
88
89     # step the environment and get new measurements
90     observation, reward, done, info = env.step(action)
91     reward_sum += reward
92
93     drs.append(reward) # record reward (has to be done after we call step() to get reward for previous action)
94
95     if done: # an episode finished
96         episode_number += 1
97
98         # stack together all inputs, hidden states, action gradients, and rewards for this episode
99         epx = np.vstack(xs)
100         eph = np.vstack(hs)
101         epdlogp = np.vstack(dlogps)
102         epr = np.vstack(drs)
103         xs,hs,dlogps,drs = [],[],[],[] # reset array memory
104
105         # compute the discounted reward backwards through time
106         discounted_epr = discount_rewards(epr)
107         # standardize the rewards to be unit normal (helps control the gradient estimator variance)
108         discounted_epr -= np.mean(discounted_epr)
109         discounted_epr /= np.std(discounted_epr)
110
111         epdlogp *= discounted_epr # modulate the gradient with advantage (PG magic happens right here.)
112         grad = policy_backward(eph, epdlogp)
113         for k in model: grad_buffer[k] += grad[k] # accumulate grad over batch
114
115         # perform rmsprop parameter update every batch_size episodes
116         if episode_number % batch_size == 0:
117             for k,v in model.iteritems():
118                 g = grad_buffer[k] # gradient
119                 rmsprop_cache[k] = decay_rate * rmsprop_cache[k] + (1 - decay_rate) * g**2
120                 model[k] += learning_rate * g / (np.sqrt(rmsprop_cache[k]) + 1e-5)
121                 grad_buffer[k] = np.zeros_like(v) # reset batch gradient buffer
122
123         # boring book-keeping
124         running_reward = reward_sum if running_reward is None else running_reward * 0.99 + reward_sum * 0.01
125         print 'resetting env. episode reward total was %f. running mean: %f' % (reward_sum, running_reward)
126         if episode_number % 100 == 0: pickle.dump(model, open('save.p', 'wb'))
127         reward_sum = 0
128         observation = env.reset() # reset env
129         prev_x = None
130
131     if reward != 0: # Pong has either +1 or -1 reward exactly when game ends.
132         print ('ep %d: game finished, reward: %f' % (episode_number, reward)) + (' ' if reward == -1 else '!!!!!!!')

```

```
if done: # an episode finished
```

```
episode_number += 1
```

```
# stack together all inputs, hidden states, action gradients, and rewards for this episode
```

```
epx = np.vstack(xs)
```

```
eph = np.vstack(hs)
```

```
epdlogp = np.vstack(dlogps)
```

```
epr = np.vstack(drs)
```

```
xs,hs,dlogps,drs = [],[],[],[] # reset array memory
```

Once a rollout is done,
Concatenate together all images, hidden
states, etc. that were seen in this batch.

Again, if using PyTorch, no need to do this.

```

64 env = gym.make("Pong-v0")
65 observation = env.reset()
66 prev_x = None # used in computing the difference frame
67 xs,hs,dlogps,drs = [],[],[],[]
68 running_reward = None
69 reward_sum = 0
70 episode_number = 0
71 while True:
72     if render: env.render()
73
74     # preprocess the observation, set input to network to be difference image
75     cur_x = prepro(observation)
76     x = cur_x - prev_x if prev_x is not None else np.zeros(D)
77     prev_x = cur_x
78
79     # forward the policy network and sample an action from the returned probability
80     aprobs, h = policy_forward(x)
81     action = 2 if np.random.uniform() < aprobs else 3 # roll the dice!
82
83     # record various intermediates (needed later for backprop)
84     xs.append(x) # observation
85     hs.append(h) # hidden state
86     y = 1 if action == 2 else 0 # a "fake label"
87     dlogps.append(y - aprobs) # grad that encourages the action that was taken to be taken (see http://cs231n.github.io/neural-networks-2/#loss)
88
89     # step the environment and get new measurements
90     observation, reward, done, info = env.step(action)
91     reward_sum += reward
92
93     ddrs.append(reward) # record reward (has to be done after we call step() to get reward for previous action)
94
95     if done: # an episode finished
96         episode_number += 1
97
98         # stack together all inputs, hidden states, action gradients, and rewards for this episode
99         eps = np.vstack(xs)
100         ephs = np.vstack(hs)
101         epdlogps = np.vstack(dlogps)
102         epr = np.vstack(drs)
103         xs,hs,dlogps,drs = [],[],[],[] # reset array memory
104
105         # compute the discounted reward backwards through time
106         discounted_epr = discount_rewards(epr)
107         # standardize the rewards to be unit normal (helps control the gradient estimator variance)
108         discounted_epr -= np.mean(discounted_epr)
109         discounted_epr /= np.std(discounted_epr)
110
111         epdlogps *= discounted_epr # modulate the gradient with advantage (PG magic happens right here.)
112         grad = policy_backward(eph, epdlogps)
113         for k in model: grad_buffer[k] += grad[k] # accumulate grad over batch
114
115         # perform rmsprop parameter update every batch_size episodes
116         if episode_number % batch_size == 0:
117             for k,v in model.iteritems():
118                 g = grad_buffer[k] # gradient
119                 rmsprop_cache[k] = decay_rate * rmsprop_cache[k] + (1 - decay_rate) * g**2
120                 model[k] += learning_rate * g / (np.sqrt(rmsprop_cache[k]) + 1e-5)
121                 grad_buffer[k] = np.zeros_like(v) # reset batch gradient buffer
122
123         # boring book-keeping
124         running_reward = reward_sum if running_reward is None else running_reward * 0.99 + reward_sum * 0.01
125         print 'resetting env. episode reward total was %f. running mean: %f' % (reward_sum, running_reward)
126         if episode_number % 100 == 0: pickle.dump(model, open('save.p', 'wb'))
127         reward_sum = 0
128         observation = env.reset() # reset env
129         prev_x = None
130
131     if reward != 0: # Pong has either +1 or -1 reward exactly when game ends.
132         print ('ep %d: game finished, reward: %f' % (episode_number, reward)) + (' ' if reward == -1 else '!!!!!!!')

```

```

# compute the discounted reward backwards through time
discounted_epr = discount_rewards(epr)
# standardize the rewards to be unit normal (helps control the gradient estimator variance)
discounted_epr -= np.mean(discounted_epr)
discounted_epr /= np.std(discounted_epr)

```

```


def discount_rewards(r):
    """ take 1D float array of rewards and compute discounted reward """
    discounted_r = np.zeros_like(r)
    running_add = 0
    for t in reversed(xrange(0, r.size)):
        if r[t] != 0: running_add = 0 # reset the sum, since this was a game boundary (pong specific!)
        running_add = running_add * gamma + r[t]
        discounted_r[t] = running_add
    return discounted_r

```

Discounted rewards

$$\sum_i A_i * \log p(y_i | x_i)$$

0.21 0.24 0.27 -0.81 -0.9 -1 0 0



Reward +1.0

Reward -1.0

```

64 env = gym.make("Pong-v0")
65 observation = env.reset()
66 prev_x = None # used in computing the difference frame
67 xs,hs,dlogps,drs = [],[],[],[]
68 running_reward = None
69 reward_sum = 0
70 episode_number = 0
71 while True:
72     if render: env.render()
73
74     # preprocess the observation, set input to network to be difference image
75     cur_x = prepro(observation)
76     x = cur_x - prev_x if prev_x is not None else np.zeros(D)
77     prev_x = cur_x
78
79     # forward the policy network and sample an action from the returned probability
80     aprobs, h = policy_forward(x)
81     action = 2 if np.random.uniform() < aprobs else 3 # roll the dice!
82
83     # record various intermediates (needed later for backprop)
84     xs.append(x) # observation
85     hs.append(h) # hidden state
86     y = 1 if action == 2 else 0 # a "fake label"
87     dlogps.append(y - aprobs) # grad that encourages the action that was taken (see http://cs231n.github.io/neural-networks-2/#loss)
88
89     # step the environment and get new measurements
90     observation, reward, done, info = env.step(action)
91     reward_sum += reward
92
93     drs.append(reward) # record reward (has to be done after we call step() to get reward for previous action)
94
95     if done: # an episode finished
96         episode_number += 1
97
98         # stack together all inputs, hidden states, action gradients, and rewards for this episode
99         eps = np.vstack(xs)
100         eph = np.vstack(hs)
101         epdlogps = np.vstack(dlogps)
102         epr = np.vstack(drs)
103         xs,hs,dlogps,drs = [],[],[],[] # reset array memory
104
105         # compute the discounted reward backwards through time
106         discounted_epr = discount_rewards(epr)
107         # standardize the rewards to be unit normal (helps control the gradient estimator variance)
108         discounted_epr -= np.mean(discounted_epr)
109         discounted_epr /= np.std(discounted_epr)
110
111         epdlogps *= discounted_epr # modulate the gradient with advantage (PG magic happens right here.)
112         grad = policy_backward(eph, epdlogps)
113         for k in model: grad_buffer[k] += grad[k] # accumulate grad over batch
114
115     # perform rmsprop parameter update every batch_size episodes
116     if episode_number % batch_size == 0:
117         for k,v in model.iteritems():
118             g = grad_buffer[k] # gradient
119             rmsprop_cache[k] = decay_rate * rmsprop_cache[k] + (1 - decay_rate) * g**2
120             model[k] += learning_rate * g / (np.sqrt(rmsprop_cache[k]) + 1e-5)
121             grad_buffer[k] = np.zeros_like(v) # reset batch gradient buffer
122
123     # boring book-keeping
124     running_reward = reward_sum if running_reward is None else running_reward * 0.99 + reward_sum * 0.01
125     print 'resetting env. episode reward total was %f. running mean: %f' % (reward_sum, running_reward)
126     if episode_number % 100 == 0: pickle.dump(model, open('save.p', 'wb'))
127     reward_sum = 0
128     observation = env.reset() # reset env
129     prev_x = None
130
131     if reward != 0: # Pong has either +1 or -1 reward exactly when game ends.
132         print ('ep %d: game finished, reward: %f' % (episode_number, reward)) + (' ' if reward == -1 else '!!!!!!!')

```

```

epdlogps *= discounted_epr # modulate the gradient with advantage (PG magic happens right here.)
grad = policy_backward(eph, epdlogps)
for k in model: grad_buffer[k] += grad[k] # accumulate grad over batch

```

$$\sum_i A_i * \log p(y_i | x_i)$$

Advantage modulation

```

def policy_backward(eph, epdlogps):
    """ backward pass. (eph is array of intermediate hidden states) """
    dw2 = np.dot(eph.T, epdlogps).ravel()
    dh = np.outer(epdlogps, model['W2'])
    dh[eph <= 0] = 0 # backpro prelu
    dw1 = np.dot(dh.T, epx)
    return {'W1':dw1, 'W2':dw2}

```

backprop!!!!!!1

```

64 env = gym.make("Pong-v0")
65 observation = env.reset()
66 prev_x = None # used in computing the difference frame
67 xs,hs,dlogps,drs = [],[],[],[]
68 running_reward = None
69 reward_sum = 0
70 episode_number = 0
71 while True:
72     if render: env.render()
73
74     # preprocess the observation, set input to network to be difference image
75     cur_x = prepro(observation)
76     x = cur_x - prev_x if prev_x is not None else np.zeros(D)
77     prev_x = cur_x
78
79     # forward the policy network and sample an action from the returned probability
80     aprobs, h = policy_forward(x)
81     action = 2 if np.random.uniform() < aprobs else 3 # roll the dice!
82
83     # record various intermediates (needed later for backprop)
84     xs.append(x) # observation
85     hs.append(h) # hidden state
86     y = 1 if action == 2 else 0 # a "fake label"
87     dlogps.append(y - aprobs) # grad that encourages the action that was taken to be taken (see http://cs231n.github.io/neural-networks-2/#loss)
88
89     # step the environment and get new measurements
90     observation, reward, done, info = env.step(action)
91     reward_sum += reward
92
93     drs.append(reward) # record reward (has to be done after we call step() to get reward for previous action)
94
95     if done: # an episode finished
96         episode_number += 1
97
98         # stack together all inputs, hidden states, action gradients, and rewards for this episode
99         eps = np.vstack(xs)
100         eph = np.vstack(hs)
101         epdlogps = np.vstack(dlogps)
102         epr = np.vstack(drs)
103         xs,hs,dlogps,drs = [],[],[],[] # reset array memory
104
105         # compute the discounted reward backwards through time
106         discounted_epr = discount_rewards(epr)
107         # standardize the rewards to be unit normal (helps control the gradient estimator variance)
108         discounted_epr -= np.mean(discounted_epr)
109         discounted_epr /= np.std(discounted_epr)
110
111         epdlogps *= discounted_epr # modulate the gradient with advantage (PG magic happens right here.)
112         grad = policy_backward(eph, epdlogps)
113         for k in model: grad_buffer[k] += grad[k] # accumulate grad over batch
114
115     # perform rmsprop parameter update every batch_size episodes
116     if episode_number % batch_size == 0:
117         for k,v in model.iteritems():
118             g = grad_buffer[k] # gradient
119             rmsprop_cache[k] = decay_rate * rmsprop_cache[k] + (1 - decay_rate) * g**2
120             model[k] += learning_rate * g / (np.sqrt(rmsprop_cache[k]) + 1e-5)
121             grad_buffer[k] = np.zeros_like(v) # reset batch gradient buffer
122
123     # boring book-keeping
124     running_reward = reward_sum if running_reward is None else running_reward * 0.99 + reward_sum * 0.01
125     print 'resetting env. episode reward total was %f. running mean: %f' % (reward_sum, running_reward)
126     if episode_number % 100 == 0: pickle.dump(model, open('save.p', 'wb'))
127     reward_sum = 0
128     observation = env.reset() # reset env
129     prev_x = None
130
131     if reward != 0: # Pong has either +1 or -1 reward exactly when game ends.
132         print ('ep %d: game finished, reward: %f' % (episode_number, reward)) + (' ' if reward == -1 else '!!!!!!!')

```

```

# perform rmsprop parameter update every batch_size episodes
if episode_number % batch_size == 0:
    for k,v in model.iteritems():
        g = grad_buffer[k] # gradient
        rmsprop_cache[k] = decay_rate * rmsprop_cache[k] + (1 - decay_rate) * g**2
        model[k] += learning_rate * g / (np.sqrt(rmsprop_cache[k]) + 1e-5)
        grad_buffer[k] = np.zeros_like(v) # reset batch gradient buffer

```

Use RMSProp for the parameter update.

RMSProp

Update rule:

$$R_t = \gamma R_{t-1} + (1 - \gamma) \nabla L_t(W_{t-1})^2$$

$$W_t = W_{t-1} - \alpha \frac{\nabla L_t(W_{t-1})}{\sqrt{R_t}}$$

Similar to AdaGrad but with an exponential moving average controlled by $\gamma \in [0, 1]$ (smaller $\gamma \implies$ more emphasis on recent gradients).


```

64 env = gym.make("Pong-v0")
65 observation = env.reset()
66 prev_x = None # used in computing the difference frame
67 xs,hs,dlogps,drs = [],[],[],[]
68 running_reward = None
69 reward_sum = 0
70 episode_number = 0
71 while True:
72     if render: env.render()
73
74     # preprocess the observation, set input to network to be difference image
75     cur_x = prepro(observation)
76     x = cur_x - prev_x if prev_x is not None else np.zeros(D)
77     prev_x = cur_x
78
79     # forward the policy network and sample an action from the returned probability
80     aprobs, h = policy_forward(x)
81     action = 2 if np.random.uniform() < aprobs else 3 # roll the dice!
82
83     # record various intermediates (needed later for backprop)
84     xs.append(x) # observation
85     hs.append(h) # hidden state
86     y = 1 if action == 2 else 0 # a "fake label"
87     dlogps.append(y - aprobs) # grad that encourages the action that was taken (see http://cs231n.github.io/neural-networks-2/#loss)
88
89     # step the environment and get new measurements
90     observation, reward, done, info = env.step(action)
91     reward_sum += reward
92
93     drs.append(reward) # record reward (has to be done after we call step() to get reward for previous action)
94
95     if done: # an episode finished
96         episode_number += 1
97
98         # stack together all inputs, hidden states, action gradients, and rewards for this episode
99         eps = np.vstack(xs)
100         eph = np.vstack(hs)
101         epdlogps = np.vstack(dlogps)
102         epr = np.vstack(drs)
103         xs,hs,dlogps,drs = [],[],[],[] # reset array memory
104
105         # compute the discounted reward backwards through time
106         discounted_epr = discount_rewards(epr)
107         # standardize the rewards to be unit normal (helps control the gradient estimator variance)
108         discounted_epr -= np.mean(discounted_epr)
109         discounted_epr /= np.std(discounted_epr)
110
111         epdlogps *= discounted_epr # modulate the gradient with advantage (PG magic happens right here.)
112         grad = policy_backward(eph, epdlogps)
113         for k in model: grad_buffer[k] += grad[k] # accumulate grad over batch
114
115         # perform rmsprop parameter update every batch_size episodes
116         if episode_number % batch_size == 0:
117             for k,v in model.iteritems():
118                 g = grad_buffer[k] # gradient
119                 rmsprop_cache[k] = decay_rate * rmsprop_cache[k] + (1 - decay_rate) * g**2
120                 model[k] += learning_rate * g / (np.sqrt(rmsprop_cache[k] + 1e-5))
121                 grad_buffer[k] = np.zeros_like(v) # reset batch gradient buffer
122
123         # boring book-keeping
124         running_reward = reward_sum if running_reward is None else running_reward * 0.99 + reward_sum * 0.01
125         print 'resetting env. episode reward total was %f. running mean: %f' % (reward_sum, running_reward)
126         if episode_number % 100 == 0: pickle.dump(model, open('save.p', 'wb'))
127         reward_sum = 0
128         observation = env.reset() # reset env
129         prev_x = None
130
131     if reward != 0: # Pong has either +1 or -1 reward exactly when game ends.
132         print ('ep %d: game finished, reward: %f' % (episode_number, reward)) + (' if reward == -1 else ' !!!!!!!!')

```

```
# boring book-keeping
```

```

running_reward = reward_sum if running_reward is None else running_reward * 0.99 + reward_sum * 0.01
print 'resetting env. episode reward total was %f. running mean: %f' % (reward_sum, running_reward)
if episode_number % 100 == 0: pickle.dump(model, open('save.p', 'wb'))
reward_sum = 0
observation = env.reset() # reset env
prev_x = None

```


```
if reward != 0: # Pong has either +1 or -1 reward exactly when game ends.
```

```
print ('ep %d: game finished, reward: %f' % (episode_number, reward)) + (' if reward == -1 else ' !!!!!!!!')
```


prints etc

In summary

1. Initialize a policy network at random
2. **Repeat Forever:**
3. Collect a bunch of rollouts with the policy
4. Increase the probability of actions that worked well
5. ???
6. Profit.



Thank you! Questions?



$$\sum_i A_i * \log p(y_i | x_i)$$

